

```
!pip install datasets transformers accelerate bitsandbytes peft trl
```

```
Requirement already satisfied: datasets in /usr/local/lib/python3.12/dist-packages (4.0.0)
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages (4.57.1)
Requirement already satisfied: accelerate in /usr/local/lib/python3.12/dist-packages (1.11.0)
Collecting bitsandbytes
  Downloading bitsandbytes-0.48.2-py3-none-manylinux_2_24_x86_64.whl.metadata (10 kB)
Requirement already satisfied: peft in /usr/local/lib/python3.12/dist-packages (0.18.0)
Collecting trl
  Downloading trl-0.25.1-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from datasets) (3.20.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages (from datasets) (2.0.2)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.12/dist-packages (from datasets) (18.1.0)
Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.12/dist-packages (from datasets) (0.3.8)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (from datasets) (2.2.2)
Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.12/dist-packages (from datasets) (2.32.4)
Requirement already satisfied: tqdm>=4.66.3 in /usr/local/lib/python3.12/dist-packages (from datasets) (4.67.1)
Requirement already satisfied: xxhash in /usr/local/lib/python3.12/dist-packages (from datasets) (3.6.0)
Requirement already satisfied: multiprocess<0.70.17 in /usr/local/lib/python3.12/dist-packages (from datasets) (0.70.16)
Requirement already satisfied: fsspec<=2025.3.0,>=2023.1.0 in /usr/local/lib/python3.12/dist-packages (from fsspec[http]=<
Requirement already satisfied: huggingface-hub>=0.24.0 in /usr/local/lib/python3.12/dist-packages (from datasets) (0.36.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.12/dist-packages (from datasets) (25.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from datasets) (6.0.3)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/dist-packages (from transformers) (2024.11.6
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages (from transformers)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.7.0)
Requirement already satisfied: psutil in /usr/local/lib/python3.12/dist-packages (from accelerate) (5.9.5)
Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from accelerate) (2.9.0+cu126)
Requirement already satisfied: aiohttp!=4.0.0a0,!=4.0.0a1 in /usr/local/lib/python3.12/dist-packages (from fsspec[http]=<
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.24
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2-
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2->datasets) (
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2->datas
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages (from requests>=2.32.2->datas
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (75.2
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (1
Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (3.1.6)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0-
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0-
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->acc
Requirement already satisfied: nvidia-nvshmem-cu12==3.3.20 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->ac
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0-
Requirement already satisfied: triton==3.5.0 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (3
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas->datasets) (
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas->datasets) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas->datasets) (2025.2)
Requirement already satisfied: aiohappyeyeballs>=2.5.0 in /usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0, !

```

```
# core idea (not full script)
from datasets import load_dataset
ds = load_dataset("greengerong/leetcode")
```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
README.md: 100%                                         21.0/21.0 [00:00<00:00, 1.81kB/s]
leetcode-train.jsonl: 100%                                     16.1M/16.1M [00:00<00:00, 17.5MB/s]
Generating train split: 100%                                    2360/2360 [00:00<00:00, 18857.27 examples/s]
```

```
ds
```

```
DatasetDict({
  train: Dataset({
    features: ['id', 'slug', 'title', 'difficulty', 'content', 'java', 'c++', 'python', 'javascript'],
```

```

        num_rows: 2360
    })
}

def make_example(example):
    title = example.get("title", "").strip()
    desc = example.get("content", "").strip()
    solution = example.get("python", "") # field names may vary; inspect dataset first
    difficulty = example.get("difficulty", "")

    # Compose instruction (short + clean)
    short_desc = "\n".join(desc.splitlines()[0:3])
    instruction = f"Solve the problem: {title}. {short_desc}..."

    # Optional input field (examples, function signature, etc.)
    input_field = example.get("content", "")

    # -----
    # Elaborated Chain-of-Thought Template
    # -----
    cot = f"""Chain-of-Thought:

1. **Problem Understanding**
The problem is titled "{title}".
The description is: {short_desc}.
We must understand the required input/output format and identify hidden constraints or edge cases.

2. **Key Observations**
- Identify crucial patterns in the problem (e.g., sorting behavior, graph structure, DP relation, prefix/suffix properties).
- Infer which constraints determine feasible complexity ( $O(n)$ ,  $O(n \log n)$ , etc.).
- Detect typical pitfalls (off-by-one, integer overflow, duplicate handling, boundary cases).

3. **Reasoning & Approach Selection**
- Consider all viable strategies.
- Discard approaches that violate time or space constraints.
- Justify the optimal algorithmic paradigm (e.g., hash map, two-pointer, BFS/DFS, DP, binary search, greedy).
- Explain why the selected approach guarantees correctness.

4. **Algorithm Breakdown**
- Provide a clear, ordered list of implementation steps.
- Handle special cases separately (empty input, zero values, duplicates, etc.).
- Present complexity analysis (time + space).

5. **Dry Run / Mini Simulation**
- Walk through a representative or tricky example step-by-step.
- Show how the algorithm transitions through states and reaches the correct output.
- Confirm correctness and edge-case behavior.

6. **Final Implementation Notes**
- Mention potential pitfalls to avoid in the final code.
- Explain any important initialization or boundary conditions.
- Ensure final code is concise, readable, and correct.

"""

    # Final answer / code
    final = solution if solution else "<write solution here>"

    output = cot + "\nFinal Answer (Python Code):\n```python\n" + final + "\n```\n\n"

    return {
        "instruction": instruction,
        "input": input_field,
        "output": output,
        "metadata": {"id": example.get("id"), "difficulty": difficulty}
    }
}

```

```

x=[]

#iterate over ds and append make_example output in X
# for i in range(len(ds["train"])):
#     x.append(make_example(ds["train"][i]))
# len(x)

for i in range(len(ds["train"])):
    x.append(make_example(ds["train"][i]))
len(x)

```

2360

```
print(x[0]["instruction"])
```

Solve the problem: Two Sum. Given an array of integers `nums` and an integer `target`, return _indices of the two numbers su

You may assume that each input would have **_exactly_ one solution**, and you may not use the _same_ element twice....

```
print(x[0]["input"])
```

Given an array of integers `nums` and an integer `target`, return _indices of the two numbers such that they add up to `targ

You may assume that each input would have **_exactly_ one solution**, and you may not use the _same_ element twice.

You can return the answer in any order.

Example 1:

Input: nums = \[2,7,11,15\], target = 9
Output: \[0,1\]

Explanation: Because nums\[0\] + nums\[1\] == 9, we return \[0, 1\].

Example 2:

Input: nums = \[3,2,4\], target = 6
Output: \[1,2\]

Example 3:

Input: nums = \[3,3\], target = 6
Output: \[0,1\]

Constraints:

* `2 <= nums.length <= 104`
* `-109 <= nums[i] <= 109`
* `-109 <= target <= 109`
* **Only one valid answer exists.**

Follow-up: Can you come up with an algorithm that is less than `O(n^2)` time complexity?

```
print(x[0]["output"])
```

Chain-of-Thought:

1. **Problem Understanding**

The problem is titled "Two Sum".

The description is: Given an array of integers `nums` and an integer `target`, return _indices of the two numbers such th

You may assume that each input would have **_exactly_ one solution**, and you may not use the _same_ element twice.. We must understand the required input/output format and identify hidden constraints or edge cases.

2. **Key Observations**

- Identify crucial patterns in the problem (e.g., sorting behavior, graph structure, DP relation, prefix/suffix properties).
- Infer which constraints determine feasible complexity ($O(n)$, $O(n \log n)$, etc.).
- Detect typical pitfalls (off-by-one, integer overflow, duplicate handling, boundary cases).

3. **Reasoning & Approach Selection**

- Consider all viable strategies.
- Discard approaches that violate time or space constraints.
- Justify the optimal algorithmic paradigm (e.g., hash map, two-pointer, BFS/DFS, DP, binary search, greedy).
- Explain why the selected approach guarantees correctness.

4. **Algorithm Breakdown**

- Provide a clear, ordered list of implementation steps.
- Handle special cases separately (empty input, zero values, duplicates, etc.).
- Present complexity analysis (time + space).

5. **Dry Run / Mini Simulation**

- Walk through a representative or tricky example step-by-step.
- Show how the algorithm transitions through states and reaches the correct output.
- Confirm correctness and edge-case behavior.

6. **Final Implementation Notes**

- Mention potential pitfalls to avoid in the final code.
- Explain any important initialization or boundary conditions.
- Ensure final code is concise, readable, and correct.

Final Answer (Python Code):

```
```python
```

```
```python
def twoSum(nums, target):
    map = {}
    for i, num in enumerate(nums):
        complement = target - num
        if complement in map:
            return [map[complement], i]
```

```
    map[num] = i  
return []
```

The algorithm leverages a hash map (unordered_map in C++, HashMap in Java, dictionary in Python, and Map in JavaScript).

This approach has a time complexity of $O(n)$ and a space complexity of $O(n)$ as well.

10

!hf auth login

```
To log in, `huggingface_hub` requires a token generated from https://huggingface.co/settings/tokens.  
Enter your token (input will not be visible):  
Add token as git credential? (Y/n) n  
Token is valid (permission: fineGrained).  
The token `varad_learning` has been saved to /root/.cache/huggingface/stored_tokens  
Your token has been saved to /root/.cache/huggingface/token  
Login successful.  
The current active token is: `varad learning`
```

```
!pip uninstall -y bitsandbytes # Force uninstall bitsandbytes  
!pip install -U bitsandbytes # Reinstall the latest bitsandbytes
```

WARNING: Skipping bitsandbytes as it is not installed.

```
Collecting bitsandbytes
  Downloading bitsandbytes-0.48.2-py3-none-manylinux_2_24_x86_64.whl.metadata (10 kB)
Requirement already satisfied: torch<3,>=2.3 in /usr/local/lib/python3.12/dist-packages (from bitsandbytes) (2.9.0+cu126)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages (from bitsandbytes) (2.0.2)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.12/dist-packages (from bitsandbytes) (25.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (3.20)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3-bit)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (75)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (1.13.3)
Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (3.1.6)
Requirement already satisfied: fsspec>=0.8.5 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (0.8.5)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-nvshmem-cu12==3.3.20 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: triton==3.5.0 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch<3,>=2.3->bitsandbytes)
Requirement already satisfied: MarkupSafe==2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2->torch<3,>=2.3->bitsandbytes)
  Downloading bitsandbytes-0.48.2-py3-none-manylinux_2_24_x86_64.whl (59.4 MB)

```

```
import accelerate
!pip install -U bitsandbytes
from transformers import AutoTokenizer, AutoModelForCausalLM, BitsAndBytesConfig # Import BitsAndBytesConfig
import torch # Import torch for dtype

model_name = "google/gemma-2-2b-bit"    # or gemma-7b-it

tokenizer = AutoTokenizer.from_pretrained(
    model_name,
    trust_remote_code=True
)

# Create a BitsAndBytesConfig object
quantization_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4", # Or "fp4"
    bnb_4bit_compute_dtype=torch.bfloat16,
    bnb_4bit_use_double_quant=True,
```

```

)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    quantization_config=quantization_config, # Pass the quantization_config object
    device_map="auto",
    trust_remote_code=True
)

Requirement already satisfied: bitsandbytes in /usr/local/lib/python3.12/dist-packages (0.48.2)
Requirement already satisfied: torch<3,>=2.3 in /usr/local/lib/python3.12/dist-packages (from bitsandbytes) (2.9.0+cu126)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages (from bitsandbytes) (2.0.2)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.12/dist-packages (from bitsandbytes) (25.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (3.20)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (75)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes))
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (3.1.6)
Requirement already satisfied: fsspec>=0.8.5 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-nvshmem-cu12==3.3.20 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bit)
Requirement already satisfied: triton==3.5.0 in /usr/local/lib/python3.12/dist-packages (from torch<3,>=2.3->bitsandbytes) (Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch<3,>=2.3->bit)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2->torch<3,>=2.3->bitsandbytes) (Requirement already satisfied: model.safetensors.index.json: 100% 24.2k/24.2k [00:00<00:00, 841kB/s]

Fetching 2 files: 100% 2/2 [00:42<00:00, 42.34s/it]
model-00002-of-00002.safetensors: 100% 241M/241M [00:17<00:00, 16.4MB/s]
model-00001-of-00002.safetensors: 100% 4.99G/4.99G [00:42<00:00, 176MB/s]
Loading checkpoint shards: 100% 2/2 [03:21<00:00, 84.91s/it]
generation_config.json: 100% 187/187 [00:00<00:00, 26.4kB/s]

```

```

tokenizer.model_max_length = 1536
tokenizer.padding_side = "right"
tokenizer.truncation_side = "right"

```

```

from peft import LoraConfig, get_peft_model

lora_config = LoraConfig(
    r=8,                                     # very important for T4
    lora_alpha=16,
    lora_dropout=0.05,
    target_modules=["q_proj", "k_proj", "v_proj", "o_proj"],
    bias="none",
    task_type="CAUSAL_LM",
)

model = get_peft_model(model, lora_config)
gradient_checkpointing = True
model.gradient_checkpointing_enable()
# Enable gradient calculation for model inputs, crucial for gradient checkpointing with PEFT
model.enable_input_require_grads()
model.print_trainable_parameters()

trainable params: 3,194,880 || all params: 2,617,536,768 || trainable%: 0.1221

```

```

from transformers import TrainingArguments

training_args = TrainingArguments(
    output_dir="gemma-leetcode-t4",
    per_device_train_batch_size=1,
    gradient_accumulation_steps=1,
    learning_rate=3e-4,

```

```

warmup_ratio=0.03,
num_train_epochs=1,
logging_steps=20,
save_steps=5000,
fp16=True,
optim="paged_adamw_8bit",
report_to="none",
remove_unused_columns=False,
gradient_checkpointing=True, # important for T4
)

# define train_dataset and data_collator
from datasets import Dataset
from transformers import DataCollatorWithPadding

def preprocess_function(examples):
    # examples contains lists under keys "instruction", "input", "output"
    inputs = []
    labels = []

    # build user-only prompts and full chat prompts (vectorized per batch)
    user_prompts = []
    full_prompts = []
    for user_text, inp_text, assistant_text in zip(examples["instruction"],
                                                   examples["input"],
                                                   examples["output"]):
        user = user_text
        if inp_text:
            user = user + "\n\nInput Examples:\n" + inp_text
        # full chat: user + assistant
        chat = [
            {"role": "user", "content": user},
            {"role": "assistant", "content": assistant_text},
        ]

        # make the string prompt using Gemma chat template (no tokenization yet)
        full_prompt = tokenizer.apply_chat_template(chat, tokenize=False)
        # user-only prompt (so we can count tokens to mask labels)
        user_chat = [{"role": "user", "content": user}]
        user_prompt = tokenizer.apply_chat_template(user_chat, tokenize=False)

        full_prompts.append(full_prompt)
        user_prompts.append(user_prompt)

    # tokenize both lists in batch
    tokenized_full = tokenizer(
        full_prompts,
        truncation=True,
        max_length=1536,
        padding=False,
        return_tensors=None
    )

    tokenized_user = tokenizer(
        user_prompts,
        truncation=True,
        max_length=1536,
        padding=False,
        return_tensors=None
    )

    # build labels: copy input_ids, then mask user tokens with -100
    out = {}
    input_ids_list = tokenized_full["input_ids"]
    attention_mask_list = tokenized_full["attention_mask"]
    user_len_list = [len(u) for u in tokenized_user["input_ids"]]

    labels_list = []
    for idx, ids in enumerate(input_ids_list):
        # copy
        lab = ids.copy()
        # mask user tokens (first user_len_list[idx] tokens) to -100
        user_len = user_len_list[idx]
        # safety: ensure we don't exceed
        if user_len > len(lab):
            user_len = len(lab)
        for j in range(user_len):
            lab[j] = -100
        labels_list.append(lab)

    out["input_ids"] = input_ids_list

```

```

out["attention_mask"] = attention_mask_list
out["labels"] = labels_list
return out

# Create the train_dataset (rehash your raw_train_dataset)
raw_train_dataset = Dataset.from_list(x)
train_dataset = raw_train_dataset.map(
    preprocess_function,
    batched=True,
    remove_columns=raw_train_dataset.column_names,
)

# use a simple padding collator (it will also pad labels properly)
data_collator = DataCollatorWithPadding(tokenizer, pad_to_multiple_of=None)

```

Map: 100% 100/100 [00:00<00:00, 126.20 examples/s]

```

from transformers import Trainer

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    data_collator=data_collator,
)

```

trainer.train()

[100/100 03:26, Epoch 1/1]

Step	Training Loss
20	0.259000
40	0.263800
60	0.258200
80	0.216900
100	0.244900

```

TrainOutput(global_step=100, training_loss=0.24854284286499023, metrics={'train_runtime': 208.1535, 'train_samples_per_second': 0.48, 'train_steps_per_second': 0.48, 'total_flos': 1404243812371968.0, 'train_loss': 0.24854284286499023, 'epoch': 1.0})

```

```

from peft import PeftModel

# Save LoRA adapters only
trainer.model.save_pretrained("gemma-2b-qlora-cot-adapter")
tokenizer.save_pretrained("gemma-2b-qlora-cot-adapter")

```

```

('gemma-2b-qlora-cot-adapter/tokenizer_config.json',
 'gemma-2b-qlora-cot-adapter/special_tokens_map.json',
 'gemma-2b-qlora-cot-adapter/chat_template.jinja',
 'gemma-2b-qlora-cot-adapter/tokenizer.model',
 'gemma-2b-qlora-cot-adapter/added_tokens.json',
 'gemma-2b-qlora-cot-adapter/tokenizer.json')

```

```

from google.colab import files
uploaded = files.upload()

```

Choose Files | No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving gemma-2b-cot-leetcode.zip to gemma-2b-cot-leetcode.zip

!unzip gemma-2b-cot-leetcode.zip

```

Archive:  gemma-2b-cot-leetcode.zip
  creating: gemma-2b-cot-leetcode/
  inflating: gemma-2b-cot-leetcode/adapter_config.json
  inflating: gemma-2b-cot-leetcode/adapter_model.safetensors
  inflating: gemma-2b-cot-leetcode/chat_template.jinja
  inflating: gemma-2b-cot-leetcode/README.md
  inflating: gemma-2b-cot-leetcode/special_tokens_map.json
  inflating: gemma-2b-cot-leetcode/tokenizer.json
  inflating: gemma-2b-cot-leetcode/tokenizer.model
  inflating: gemma-2b-cot-leetcode/tokenizer_config.json

```

```

from transformers import AutoModelForCausalLM, AutoTokenizer
from peft import PeftModel

```

```

base_model_name = "google/gemma-2-2b-it"    # Changed to match the likely base model of the adapter
lora_path = "gemma-2b-cot-leetcode"

tokenizer = AutoTokenizer.from_pretrained(base_model_name)
base_model = AutoModelForCausalLM.from_pretrained(
    base_model_name,
    torch_dtype="auto",
    device_map="auto"
)

model = PeftModel.from_pretrained(base_model, lora_path)
model.eval()

tokenizer_config.json: 100%                                47.0k/47.0k [00:00<00:00, 5.46MB/s]
tokenizer.model: 100%                                    4.24M/4.24M [00:00<00:00, 229kB/s]
tokenizer.json: 100%                                    17.5M/17.5M [00:00<00:00, 34.3MB/s]
special_tokens_map.json: 100%                            636/636 [00:00<00:00, 44.4kB/s]
config.json: 100%                                     838/838 [00:00<00:00, 107kB/s]
model.safetensors.index.json: 100%                      24.2k/24.2k [00:00<00:00, 2.63MB/s]
Fetching 2 files: 100%                                2/2 [01:24<00:00, 84.47s/it]
model-00001-of-00002.safetensors: 100%                  4.99G/4.99G [01:24<00:00, 64.5MB/s]
model-00002-of-00002.safetensors: 100%                  241M/241M [00:36<00:00, 6.02MB/s]
Loading checkpoint shards: 100%                        2/2 [00:23<00:00, 9.81s/it]
generation_config.json: 100%                          187/187 [00:00<00:00, 21.2kB/s]

PeftModelForCausalLM(
    (base_model): LoraModel(
        (model): Gemma2ForCausalLM(
            (model): Gemma2Model(
                (embed_tokens): Embedding(256000, 2304, padding_idx=0)
                (layers): ModuleList(
                    (0-25): 26 x Gemma2DecoderLayer(
                        (self_attn): Gemma2Attention(
                            (q_proj): lora.Linear(
                                (base_layer): Linear(in_features=2304, out_features=2048, bias=False)
                                (lora_dropout): ModuleDict(
                                    (default): Dropout(p=0.05, inplace=False)
                                )
                            )
                            (lora_A): ModuleDict(
                                (default): Linear(in_features=2304, out_features=8, bias=False)
                            )
                            (lora_B): ModuleDict(
                                (default): Linear(in_features=8, out_features=2048, bias=False)
                            )
                            (lora_embedding_A): ParameterDict()
                            (lora_embedding_B): ParameterDict()
                            (lora_magnitude_vector): ModuleDict()
                        )
                        (k_proj): lora.Linear(
                            (base_layer): Linear(in_features=2304, out_features=1024, bias=False)
                            (lora_dropout): ModuleDict(
                                (default): Dropout(p=0.05, inplace=False)
                            )
                            (lora_A): ModuleDict(
                                (default): Linear(in_features=2304, out_features=8, bias=False)
                            )
                            (lora_B): ModuleDict(
                                (default): Linear(in_features=8, out_features=1024, bias=False)
                            )
                            (lora_embedding_A): ParameterDict()
                            (lora_embedding_B): ParameterDict()
                            (lora_magnitude_vector): ModuleDict()
                        )
                    )
                )
            )
        )
    )
)

def generate(model, tokenizer, instruction, input_text):
    user = instruction
    if input_text:
        user = user + "\n\nInput:\n" + input_text

        # Provide a steering / instruction sentence (safe, short) in the user prompt
        steering = " Please provide a clear step-by-step reasoning labeled 'Step 1:', 'Step 2:', ... followed by the final answer"
        user = user + steering

    messages = [{"role": "user", "content": user}]
    inputs = tokenizer.apply_chat_template(
        messages,
        add_generation_prompt=True,
        tokenize=True,
        return_dict=True,
    )

```

```

        return_tensors="pt",
    ).to(model.device)

outputs = model.generate(**inputs, max_new_tokens=5000)
return tokenizer.decode(outputs[0][inputs["input_ids"].shape[-1]:])

        (default): Linear(in_features=8, out_features=2304, bias=False)

# -----
# Inference on a random example
# -----
import random

idx = random.randint(0, len(x)-1)
test = x[idx]

print("\n★ TEST EXAMPLE:")
print(test["instruction"])

result = generate(
    model,
    tokenizer,
    instruction=test["instruction"],
    input_text=test["input"]
)

print("\n★ MODEL OUTPUT:")
print(result)

print("\n★ EXPECTED OUTPUT:")
print(test["output"])

```

★ TEST EXAMPLE:
Solve the problem: Spiral Matrix II. Given a positive integer `n`, generate an `n x n` `matrix` filled with elements from

Example 1:...

★ MODEL OUTPUT:
Solution:

Step 1: We will use a 2D array to store the matrix. The size of the matrix will be `n x n`.

Step 2: We will iterate through the matrix in a spiral order. The spiral order is defined as follows:

- * Start at the top-left corner of the matrix.
- * Move right until you reach the right edge of the matrix.
- * Move down until you reach the bottom edge of the matrix.
- * Move left until you reach the left edge of the matrix.
- * Move up until you reach the top edge of the matrix.
- * Repeat the process.

Step 3: We will use a counter to keep track of the current element in the matrix.

Step 4: We will use a nested loop to iterate through the matrix.

Step 5: We will increment the counter and add the current element to the matrix.

Step 6: We will decrement the counter and add the current element to the matrix.

Step 7: We will repeat the process until we have filled the entire matrix.

Step 8: We will return the matrix.

Code:

```

```python
def generateMatrix(n):
 matrix = [[0 for _ in range(n)] for _ in range(n)]
 counter = 1
 top, bottom = 0, n - 1
 left, right = 0, n - 1

 while top <= bottom and left <= right:
 for i in range(left, right + 1):
 matrix[top][i] = counter
 counter += 1
 top += 1
 for i in range(top, bottom + 1):
 matrix[i][right] = counter
 counter += 1
 right -= 1
 if top <= bottom and left <= right:
 for i in range(right, left - 1, -1):
 matrix[bottom][i] = counter
 counter += 1
```

```

```
bottom -= 1
for i in range(bottom, top - 1, -1):
    matrixfillleft = counter
```